

Search engine for generating a movie cast

My project subtitle

Christoffer Holmesland
University of Stavanger, Norway
cr.holmesland@stud.uis.no

Malavika Ramakrishnan
University of Stavanger, Norway
m.ramakrishnan@stud.uis.no

ABSTRACT

In this report we show how we used data from IMDB(source) to generate a movie cast based on user input. To determine the recommended cast we calculate a score based on genres, actor to actor relationship, and a brief movie summary. The user is presented with a list of actors and the predicted IMDB user rating score.

KEYWORDS

Hadoop, Spark, Film, Casting, Search engine

1 INTRODUCTION

Context and Motivation. For a low budget movie it is important to have a well-known actor. It can increase the box office sales or help attract investors (source stephenfollows). If you are new to the film industry it can be hard to know which actors or actresses will be a good fit for your movie. It is possible to use previous movies where you liked the performance of the actor as inspiration, but there are too many movies being made that it is not possible to get a proper overview of all the possibilities. Another possibility is to look at websites like IMDB (source) to find inspiration. You can use features like the user score of a movie to select actors. We do not consider this as a good option because there are too many movies to compare, and in practice you will end up looking at only the top rated movies.

Research Problem. In this report we present a solution to the problem of selecting a movie cast. Our system lets the user input a brief summary of their movie and a list of actor characteristics they want the lead actors to have. The system uses data from IMDB to rank possible actors and returns a list of suggestions to the user. The rank is based on a combination of past performance, acting relationships and a subset of the actor characteristics.

Contribution Summary.

- We implement PAXOS using brand new technologies.

2 RELATED WORK

The IMDB datasets have been used to predict the gross movie revenue (source jae huang). The accuracy of several machine learning models were compared to find the best way of predicting the gross revenue. The results show that the random forest model had the best performance and that the most important feature is the number of user votes. The datasets have also been used to compare the user rating with other features (source max woelf). In the same analysis

they found that lead actors are typically ten years older than lead actresses.

The movie project system cinelytic developed by Cinelytic, Inc. (source) is able to analyze and produce forecasts of how well a movie will perform based on country, cast and distributor. The system ranks actors based on the predicted economic increase they have for a given project.

Our system is similar to cinelytic, but we are basing the score only on IMDB data and use a subset of the features considered by cinelytic.

3 BACKGROUND

Give a short, self-contained summary of necessary background information.

For our example if your project is a new PAXOS implementation, the background would be a brief introduction to the consensus problem, the PAXOS algorithm, the system model used by the PAXOS algorithm and common optimizations and relevant variations of that algorithm.

The goal of the background section is to make the paper self-contained for an audience as large as possible. As in every section you start with a very brief overview of the section. For our example that would be: In this section, we introduce the consensus problem, the PAXOS algorithm and common optimizations of that algorithm.

If you copy a definition from a text book or some other source, you must cite that source. See Example 8.1 below.

4 HADOOP CLUSTER

Our Hadoop cluster consists of one name node and three slave nodes. The name node has 8 GB RAM and a 2.4 GHz Intel Skylake processor with 4 cores. The three slave nodes have 4 GB RAM each and a 2.4 GHz Intel Skylake processor with 2 cores. We used Hadoop version 3.2.1, Spark version 3.0.0-preview2 and Python version 3.5.2.

5 YOUR METHOD

5.1 Getting the data

First we had to get all of the required data. Most of it can be downloaded from the IMDB website(source to imdb/interfaces). This data does not include actor gender, movie plot or box office values. To determine the gender we used the list of known professions for every person. If they are known for being a actor they are assigned the label "0". Actresses are assigned the label "1". Some people in the dataset have never acted, for example directors or producers. They are removed from the data. There is no way to find the movie plot in the datasets, so we had to find an external source. One alternative is to use webscraping on the IMDB website. We tried doing this

Supervised by Leander Jehl and Hein Meling.

Project in Data-intensive Systems (DAT500), IDE, UiS
2020.

but quickly realized that it would not be possible for us because they stop responding to requests after too many in a given time period. We were able to get about 9 000 requests per hour. Since the dataset includes about 500 000 movies this would not be a good solution. It was also not possible for us to know if there were other rate limits e.g. 50 000 requests per day. Instead we used the OMDB API (source) to retrieve the movie plots. Using this API we were able to get all of the plots in 1 hour and 43 minutes. Our original idea was to also make a prediction on the box office value. After using the OMDB API to also find the box office values we saw that only 6381 movies have this value. We do not believe this is enough data to make proper predictions on (why?) so we did not use this data. Instead we decided to generate a prediction for the IMDB user score because this value is available for all of the movies in our dataset.

5.2 Preprocessing

The data from IMDB includes a lot of information we don't need and it also has missing values. It is also split over several files e.g. title.basics.tsv and title.ratings.tsv. Every person in the dataset has an id, name, birthyear, deathyear, a list of primary professions and a list of titles they are known for working on. We replace the primary profession list with a number indicating the gender. We remove people who are dead or who has never been an actor or actress before. There are also some people who are actors but not known for any titles. They are also removed because we need that information to calculate a score. These steps reduce the number of people from 9.9 million to 200 thousand.

The data in the title.basics.tsv file is an id, type, well known title, original title, adult, start year, end year, runtime, and a list of genres. The only relevant information for our project is the id, type and list of genres. The type is used to decide if we should keep it. We are only interested in looking at movies, represented by "movie" or "tvMovie" in the data. Shorts are typically news segments which would not be a good indicator of movie performance. If we included tv series then the actors in them would have inflated scores because there could be over 100 episodes for a given series. That would show that the person is good at that specific role, but not represent how well he or she can adapt to a new role. The runtime is not used in our system because it is missing for most of the titles. Some movies are missing the genre list and are thus removed because it is required for the system. This reduces the number of titles from 6.5 million to 580 thousand. This might seem like a lot, but most of it is because the data includes separate entries for each season and episode for every tv series. In this step we also use the plot data from the OMDB API. Movies without the plot information are removed. This reduces the number further to 212 thousand (check this number again).

The movie ratings are in the title.ratings.tsv file from IMDB. The format is id, average rating and number of votes. Any title removed from title.basics.tsv is also removed from this data. The title.principals.tsv file contains a list of the most important people for each title. The fields are title id, ordering, actor id, category, job, characters. We use this information to determine what actors have worked together. We only look at the titles with id which was not removed in the previous step, and the ordering, category, job

and characters fields are dropped because they provide no relevant information to us.

To improve the runtime performance of the search we calculate and store several actor dependent values during preprocessing. For every actor we calculate how many unique people they have worked with, and how many people they have worked with in total. This is used in our ranking function to give a measure of how well they work with other people. The other scores we calculate are genre scores. It is calculated as the average of the movie rating the actor has been in. The score is calculated for each genre. If actor A has acted in two movies, the first with score 5 with genres "romance" and "drama", the second with score 10 and genre "romance", then actor A has a "romance" score of 7.5 and a "drama" score of 5.

6 EXPERIMENTAL EVALUATION

Here you evaluate your work using experiments. You start again with a very short summary of the section. The typical structure follows.

6.1 Experimental setup

Specify the context and setup of your experiments. This includes e.g. what hardware (VMs) you are running on, what operating system these machines are running, how they are connected, ...

Also explain how you generate load for your system and what parameters you used here. The general idea is to include enough information for others to reproduce your experiments. To that end, you should provide a detailed set of instructions for repeating your experiments. These instructions should not be included in the report, but should be provided as part of the source code repository on GitHub, typically as the README.md file, or as Shell scripts or Ansible scripts.

If your experiments give strange or unexpected results, analyze, profile and debug your code. **Do not simply re-run experiments until they give the expected results.**

Finally, running experiments is very time consuming and you may need to go through multiple rounds of experiment, debugging and optimization. **Do not delay running experiments until the end of your project period.**

6.2 Results

The results of your experiments. Compare different variants of your design (e.g. with and without optimizations) or compare performance to other designs or systems. Plots should show the average over multiple runs (at least 10 as a rule of thumb), including error bars, percentiles or min/max values.

Discuss the plot and extract the overall performance. Do not repeat all numbers in the text, but mention relevant differences in numbers, e.g. our optimization improves throughput by 26%. Discuss how the results validate or contradict your assumptions.

Perform experiments to evaluate your system under operating normal conditions, when experiencing failures or attacks, or with different workloads.

Figure 1 shows a graph generated with pgfplots from experiment data.

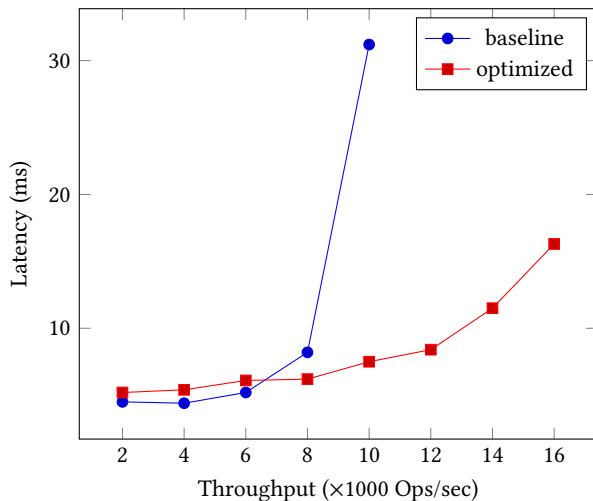


Figure 1: A graph showing latency and throughput of a baseline and optimized implementation. The axes show latency in milliseconds, and throughput in thousand operations per second. Data is made up.

7 CONCLUSION

Here you need to summarize what you did and why this is important. Do not take the abstract and put it in the past tense. Remember, now the reader has (hopefully) read the paper, so it is a very different situation from the abstract. Try to highlight important results and say the things you really want to get across such as high-level statements (e.g., we believe that is the right approach to Even though we only considered LAN, the technique should be applicable) You can also formulate next steps if you want. Be brief.

8 CITATION

It is important that you correctly refer to sources.

- If you describe background, according to some textbook, cite that textbook. *See Example 8.1 below.*
- If you make a claim about trends in industry or research, add citations to validate the claim. *See Example 8.2 below.*
- **Do not copy paste text from any source into your report without citation. It will be considered plagiarism.**
- Always use the tilde character between the text and the cite command, e.g.:

DBLP~\cite{dplg}

The tilde inserts a space, but prevents line break between the text and citation.

- For correct citations we recommend that you export bibtex entries from DBLP [4]. But make sure to update the bibtex entries to adhere to the following rules.

Here are some rules to be followed for your bibliography:

- Make sure titles are correctly capitalized.
- Make sure there is enough info for a reader to find the document: authors plus title isn't enough.
- Don't blindly copy and paste ACM bibtex entries without fixing them:

- they list "New York, NY" as the address field, which means that all ACM conferences look like they took place there; just delete this, or - better
- replace it with where the conference took place, which is much more helpful to the reader than the publisher's address
- they have a truly weird idea about how to capitalize conference names; just correct this (consistently)
- * IEEE is equally bad, and uses a partially-inverted form of conference names.
- * Please actually review the bibliography before submitting your paper!

And while you are at it, please:

- Add `includemathpmtc` in the main LaTeX file, so you don't have math stuff in a different font than the main text.
- Always use the tilde character between the number and unit, e.g. 100 Mbps or 53 ms. The tilde inserts a space, but prevents line break between the number and unit.
- Never put SI units in italics.
- Do differentiate between bits (b) and bytes (B), and between powers of 10 (MB) and 2 (MiB).
- Avoid things like: "We refer the reader to [42]." That is, don't use citations as nouns.

Example 8.1. Consensus is a distributed programming abstraction that fulfills the following properties [2]:

Termination

Validity

Integrity

Agreement

Example 8.2. Consensus systems build a fundamental part of today's cloud-computing infrastructure [1, 3, 5, 6].

9 LATEX

This section contains some instructions and examples for using \LaTeX .

9.1 You can add subsections

You can use numbered subsections to structure your sections or `\paragraph` to separate paragraphs with a heading using only a few words, as used in Section 1.

- (1) This is an item in an enumeration.
 - This is an item of a unnumbered list. In this case, the lists are nested within each other.
- (2) The second point in the enumerated list.

I can refer to the element of the enumeration above as Point 1. If you refer to numbered items, e.g. items from a list or figures or sections, always capitalize the name. For example, this is Section 9.

9.2 Figures

You can include figures. You can include files, as done in Figure 2. Avoid including jpeg, gif or bmp files since these do not scale nicely. Again Figure 2 is an example for that.

You can create graphs from your experiment-data using `pgfplots`. See the example in `tex/evaluation.tex` and documentation <http://pgfplots.sourceforge.net/pgfplots.pdf>.




Figure 2: Figure taken from [7]

9.3 Other tips

- Always use the tilde character between the number and unit, e.g. 100 Mbps or 53 ms. The tilde inserts a space, but prevents line break between the number and unit.
- Never put SI units in italics.
- Do differentiate between bits (b) and bytes (B), and between powers of 10 (MB) and 2 (MiB).
- Avoid things like: "We refer the reader to [42]." That is, don't use citations as nouns.

For further instructions on how to add **Tables**, **Algorithms**, **Theorems** see [acmguide.pdf](#).

10 FURTHER READING

Here are some other useful resources: (make these into references instead of links?)

- The co-chairs of the Ninth SOSP prepared some recommendations and valuable lessons in a paper titled *How (and How Not) to Write a Good Systems Paper* https://www.usenix.org/legacy/publications/library/proceedings/dsl97/good_paper.html
- <http://gramoli.redbellyblockchain.io/web/doc/talks/researchmethod.pdf>
- https://www.doc.ic.ac.uk/~prp/doc/talks/11-prp-paper_writing.pdf
- This template is inspired by a template from Markus Püschel [8].

REFERENCES

- [1] Michael Burrows. 2006. The Chubby Lock Service for Loosely-Coupled Distributed Systems. In *7th Symposium on Operating Systems Design and Implementation (OSDI '06), November 6-8, Seattle, WA, USA*. 335–350. <http://www.usenix.org/events/osdi06/tech/burrows.html>
- [2] Christian Cachin, Rachid Guerraoui, and Luís E. T. Rodrigues. 2011. *Introduction to Reliable and Secure Distributed Programming* (2. ed.). Springer. <https://doi.org/10.1007/978-3-642-15260-3>
- [3] James C. Corbett, Jeffrey Dean, Michael Epstein, Andrew Fikes, Christopher Frost, J. J. Furman, Sanjay Ghemawat, Andrey Gubarev, Christopher Heiser, Peter Hochschild, Wilson C. Hsieh, Sebastian Kanthak, Eugene Kogan, Hongyi Li, Alexander Lloyd, Sergey Melnik, David Mwaura, David Nagle, Sean Quinlan, Rajesh Rao, Lindsay Rolig, Yasushi Saito, Michal Szymaniak, Christopher Taylor, Ruth

- Wang, and Dale Woodford. 2013. Spanner: Google's Globally Distributed Database. *ACM Trans. Comput. Syst.* 31, 3 (2013), 8:1–8:22. <https://doi.org/10.1145/2491245>
- [4] Schloss Dagstuhl and University of Trier. 2018. DPLG Computer Science Bibliography. <http://dblp.uni-trier.de/>
- [5] Apache Software Foundation. 2010. ZooKeeper. <https://zookeeper.apache.org/>
- [6] Patrick Hunt, Mahadev Konar, Flavio Paiva Junqueira, and Benjamin Reed. 2010. ZooKeeper: Wait-free Coordination for Internet-scale Systems. In *2010 USENIX Annual Technical Conference, Boston, MA, USA, June 23-25, 2010*. <https://www.usenix.org/conference/usenix-atc-10/zookeeper-wait-free-coordination-internet-scale-systems>
- [7] Leander Jehl. 2018. Lecture Notes for DAT520 Distributed Systems.
- [8] Markus Püschel. 2011. A Descriptive Title, not too general, not too long. <https://www.inf.ethz.ch/personal/markusp/teaching/263-2300-ETH-spring11/project/report.pdf>